

## Introduction to ggplot2

STA3100: Programming with Data

# ggplot2

- ▶ Is an R package for producing statistical, or data, graphics.
- ▶ ggplot2 has an underlying grammar, based on the **G**rammar of **G**raphics, that allows you to compose graphs by combining independent components.
- ▶ You can create novel graphics that are tailored to your specific problem.
- ▶ Book: *ggplot2: Elegant Graphics for Data Analysis* by Hadley Wickham, Danielle Navarro, and Thomas Lin Pedersen (<https://ggplot2-book.org/>).

## Differences with base graphics

- ▶ `ggplot()` is the main function used to make any plot.
- ▶ The function works with data frames and not individual vectors.
- ▶ You can keep enhancing a plot by adding more layers (and themes) to an existing plot created using the `ggplot()` function.

## Example: Fuel economy data

Includes information about the fuel economy of popular car models in 1999 and 2008, collected by the US Environmental Protection Agency.

```
library(ggplot2)
mpg
```

```
## # A tibble: 234 x 11
##   manufacturer model    displ  year  cyl trans  drv    cty   hwy fl    class
##   <chr>          <chr>    <dbl> <int> <int> <chr>  <chr> <int> <int> <chr> <chr>
## 1 audi          a4        1.8  1999    4 auto(l~ f    18    29 p    comp~
## 2 audi          a4        1.8  1999    4 manual~ f    21    29 p    comp~
## 3 audi          a4        2    2008    4 manual~ f    20    31 p    comp~
## 4 audi          a4        2    2008    4 auto(a~ f    21    30 p    comp~
## 5 audi          a4        2.8  1999    6 auto(l~ f    16    26 p    comp~
## 6 audi          a4        2.8  1999    6 manual~ f    18    26 p    comp~
## 7 audi          a4        3.1  2008    6 auto(a~ f    18    27 p    comp~
## 8 audi          a4 quat~  1.8  1999    4 manual~ 4    18    26 p    comp~
## 9 audi          a4 quat~  1.8  1999    4 auto(l~ 4    16    25 p    comp~
## 10 audi         a4 quat~  2    2008    4 manual~ 4    20    28 p    comp~
## # ... with 224 more rows
```

## Example: Fuel economy data

The variables are mostly self-explanatory:

- ▶ `cty` and `hwy` record miles per gallon (mpg) for city and highway driving.
- ▶ `displ` is the engine displacement in liters (engine size).
- ▶ `drv` is the drivetrain: front wheel (f), rear wheel (r) or four wheel (4).
- ▶ `model` is the model of car. There are 38 models, selected because they had a new edition every year between 1999 and 2008.
- ▶ `class` is a categorical variable describing the “type” of car: two seater, SUV, compact, etc.

# Tibbles vs Data Frames

- ▶ Tibbles are data frames with a tweak on some older behaviours:
  1. Refined print method shows only the first 10 rows and all the columns that fit on screen. Includes column name and type.
  2. Subsetting can be done with `$` by name, and `[[` by name or position.

```
mpg$cty  
mpg[["cty"]]  
mpg[[8]]
```

- ▶ Most packages use regular data frames (`data.frame()`). Use `as.data.frame()` to turn a tibble back to a data.frame.
- ▶ Use `as_tibble()` to coerce a data.frame to a tibble.


## Key components ggplot2 plots

Every plot has three key components:

1. **data**,
2. A set of **aesthetic mappings** between variables in the data and visual properties, and
3. At least one layer which describes how to render each observation. Layers are usually created with a **geom** function.

## Example: Scatterplot of engine size vs hwy mileage

```
ggplot(mpg, aes(x = displ, y = hwy)) +  
  geom_point()
```





## Example: Scatterplot of engine size vs hwy mileage

This produces a scatterplot defined by:


1. Data: mpg.
2. Aesthetic mapping: engine size mapped to x position, fuel economy to y position.
3. Layer: points.

# Summary of ggplot2 plots

## 1. With `ggplot()`:

- ▶ The first argument is always the name of the dataset you wish to use for plotting.
- ▶ Next, you provide the variables from the dataset to be assigned to different aesthetic elements of the plot, such as the x and the y axes.

```
ggplot(mpg, aes(x = displ, y = hwy))
```



## Summary of ggplot2 plots (contd.)

2. Next, you tell `ggplot()` what type of visualization you would like to add to the blank template.
  - ▶ You add another layer to the `ggplot()` by adding a `+` at the end of the line, to indicate that you are adding a layer
  - ▶ then specify the **geometric** object to be used to create the plot.

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point()
```

This tells `ggplot()` that each data point should be represented by one point on the plot. The first two unnamed arguments to `aes()` will always be mapped to `x` and `y`.

## Other Plot geoms

You can substitute `geom_point()` for a different geom function to get a different type of plot.

- ▶ `geom_boxplot()` produces a box-and-whisker plot to summarise the distribution of a set of points. Use `geom_violin()` for violin plots.
- ▶ `geom_histogram()` and `geom_freqpoly()` show the distribution of continuous variables.
- ▶ `geom_bar()` shows the distribution of categorical variables.
- ▶ `geom_line()` draw lines between the data points.

## Colour, size, shape and other aesthetic attributes

To add additional variables to a plot, we can use other aesthetics like colour, shape, and size into the call to `aes()`:


- ▶ `aes(displ, hwy, color = class)`, better with categorical.
- ▶ `aes(displ, hwy, shape = drv)`, better with categorical.
- ▶ `aes(displ, hwy, size = cyl)`, better with numerical.

The following code gives each point a unique color corresponding to its class:

```
ggplot(mpg, aes(displ, hwy, color = class)) +  
  geom_point()
```


## Example: Scatterplot of engine size vs hwy mileage by drivetrain

```
ggplot(mpg, aes(displ, hwy, shape = drv, color = drv)) +  
  geom_point()
```



Example: Scatterplot of engine size vs hwy mileage by cylinders


```
ggplot(mpg, aes(displ, hwy, size = cyl)) +  
  geom_point()
```



## Example: Aesthetic to a fixed value

If you want to set an aesthetic to a fixed value, without scaling it, do so in the individual layer outside of `aes()`.

```
ggplot(mpg, aes(displ, hwy)) +  
  geom_point(colour = "blue", size=4)
```






## Example: Boxplots of hwy mileage by drivetrain

Boxplots summarise the bulk of the distribution with the median, IQR and whiskers.


```
ggplot(mpg, aes(drv, hwy)) +  
  geom_boxplot(fill="lightblue")
```



## Example: Violin plots of hwy mileage by drivetrain

Violin plots give the richest display, but rely on the calculation of a density estimate, which can be hard to interpret.

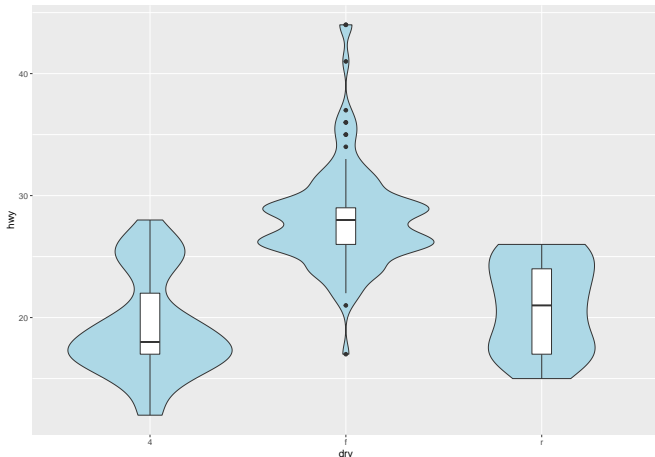
```
p <- ggplot(mpg, aes(drv, hwy)) +  
  geom_violin(fill="lightblue")  
p
```



## Example: Adding layers to plots

Violin plots of hwy mileage by drivetrain with boxplots


```
p + geom_boxplot(width=0.1)
```



## Example: Histograms of highway mpg

Histograms show the distribution of a single numeric variable.


```
ggplot(mpg, aes(hwy)) + geom_histogram()
```



## Example: Histograms of highway mpg


- ▶ You can control the width of the bins with the `binwidth` argument.
- ▶ The default just splits your data into 30 bins (`binwidth = 1`).
- ▶ Try many bin widths to tell the full story of your data.

```
ggplot(mpg, aes(hwy)) + geom_histogram(binwidth = 0.5)
```



## Example: Frequency polygons of highway mpg


```
ggplot(mpg, aes(hwy)) + geom_freqpoly(binwidth = 2.5) +  
  geom_histogram(binwidth = 2.5, color = "black",  
                 fill = "lightblue")
```



## Example: Frequency polygons of engine size by drivetrain

To compare the distributions of different subgroups, you can map a categorical variable to either fill (for `geom_histogram()`) or color (for `geom_freqpoly()`):


```
ggplot(mpg, aes(displ, color = drv)) +  
  geom_freqpoly(binwidth = 0.5)
```



## Example: Bar chart of manufacturer

The discrete analogue of the histogram is the bar chart:

```
ggplot(mpg, aes(manufacturer)) +  
  geom_bar(fill = "lightblue")
```



► Note that we are using unsummarised data.



## Example: Bar chart for summarised data

- ▶ Toy example of data frame

```
drugs <- data.frame(drug = c("a", "b", "c"),  
  effect = c(4.2, 9.7, 6.1))
```

- ▶ To display this sort of data, you need to tell `geom_bar()` to not run the default stat which bins and counts the data:

```
ggplot(drugs, aes(drug, effect)) +  
  geom_bar(stat = "identity")
```

